

K8s Security



What is SSL/TLS?

- ◆ SSL = Secure Sockets Layer
- ◆ TLS = Transport Layer Security
- SSL (Secure Sockets Layer) and its modern version TLS (Transport Layer Security) are security protocols.
- They ensure that data transferred between a client and a server is encrypted and safe from hackers.
- They are the reason why websites use `https://` instead of `http://`.
- `https://` means the site is secure and encrypted, while `http://` is not secure.

TLS is the newer and more secure version of SSL.

They are protocols used to encrypt data sent between:

- Our browser (client)
- A website (server)



Why Do We Need It?

Imagine this:



User (Client) → sends a request via HTTP → Server

Let's say we're on a public Wi-Fi (like in a coffee shop) and type our **bank password** on a website.

- If the site is **http://**, our data (password) goes in **plain text**. A hacker can **see it**.
- If the site is **https://**, our data is **encrypted** using TLS. A hacker **cannot read it**.

That's why most sites now use **https://**.

If it's just **HTTP**, the data (like passwords, bank info) travels in **plain text**.

 A hacker in between can intercept and read everything.

Symmetric Encryption (Why It's Not Safe Alone):

1 User Sends Request to Server:

- When a user wants to connect to a server (like logging into a website), the server may ask for identity verification.
- This is like showing our ID card at the entrance—it proves who we are before giving access.

2 User Sends Encrypted Data + Key:

- In symmetric encryption, both parties use the **same key** to encrypt and decrypt data.
- So the user must send their encrypted data **along with the secret key** to the server.

3 Server Uses the Key to Decrypt:

- The server uses this same key to decrypt the user's data.

4 Why It's Not Safe:

- While the data and key are being sent over the internet, a hacker can intercept both.
- This is like giving your locker key and our locker contents together—if someone catches it mid-way, they get full access.
- **Result:** Both our message and the secret key are exposed → the entire system is broken.



TLS/SSL (Asymmetric Encryption)

TLS (Transport Layer Security) uses **asymmetric encryption**, which is more secure than symmetric. It works using **two keys**:

-  **Public Key**—Shared with anyone (user/client)
-  **Private Key**—Kept secret by the server

Here's how the secure connection is built:

1 User Sends Request to the Server:

- When we visit a secure website (like `https://example.com`), our browser sends a request to connect.

2 Server Creates a CSR (Certificate Signing Request):

- The server generates a **CSR**, which includes its **public key** and identity info (like domain name, organization, location).

3 CSR is Sent to CA (Certificate Authority):

- The server sends this CSR to a trusted **Certificate Authority (CA)**.

- The CA verifies the server's identity and then **digitally signs** the certificate.

```
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl certificate deny admin
certificatesigningrequest.certificates.k8s.io/admin denied
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl get csr
NAME    AGE    SIGNERNAME           REQUESTOR      REQUESTEDDURATION   CONDITION
admin   57s    kubernetes.io/kube-apiserver-client  kubernetes-admin  24h        Denied
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl certificate approve admin
error: certificate signing request "admin" is already Denied
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ |
```

```
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl get csr
NAME    AGE    SIGNERNAME           REQUESTOR      REQUESTEDDURATION   CONDITION
admin   12s   kubernetes.io/kube-apiserver-client  kubernetes-admin  24h        Pending
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl certificate approve admin
certificatesigningrequest.certificates.k8s.io/admin approved
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl get csr
NAME    AGE    SIGNERNAME           REQUESTOR      REQUESTEDDURATION   CONDITION
admin   33s   kubernetes.io/kube-apiserver-client  kubernetes-admin  24h        Approved,Issued
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ |
```

4 Signed Certificate is Sent to the User:

- The server sends this **signed certificate** (includes the public key and CA's signature) to the user's browser.
- This tells the browser: “I'm a legit website, and here's proof from a trusted authority.”

5 User Uses Public Key for Encryption:

- Now the user encrypts a **temporary symmetric key** using the **server's public key** and sends it back.

6 Only Server Can Decrypt It:

- Since only the server has the **private key**, only it can decrypt the symmetric key.
- From now on, both user and server use this symmetric key for faster secure communication.

 **Result:** Even if a hacker sees the public key, they **can't decrypt** anything without the private key.

What About User Authentication?

Until now, we talked about **server authentication** (browser verifies server).

But sometimes, the **server also needs to verify the user's identity**.

Two Types:

1. **Public Websites (like Amazon):**

- User authentication happens **later**, like when we **log in** with username and password.



Real-Life Example (Putting it All Together)

Let's say we open `https://www.amazon.com` :

1. Our browser sends a request to Amazon.
2. Amazon sends a certificate to prove it's really Amazon.
3. Our browser checks the certificate.
4. If valid, our browser generates a session key.
5. Encrypts it using Amazon's public key.
6. Amazon decrypts it using its private key.
7. Now both share a session key.
8. All further data (login, payment, etc.) is encrypted.

Done securely. No hacker can see our password or card details.

2. Private Systems (like internal company tools or APIs):

In internal/private environments (such as enterprise applications, DevOps pipelines, or internal APIs), companies often want **stronger user authentication**—not just passwords.

 Mutual TLS (mTLS) is used for this.

What is Mutual TLS?

Normally, in regular TLS:

- The **server** proves its identity to the **client** (like our browser).
- The **client** is anonymous (it just trusts the server).

But in **Mutual TLS**:

- Both **server** and **client** must prove their identity to each other using **certificates**.

How does it work in private systems?

Custom CA (Certificate Authority):

- Companies set up their own **internal CA** to issue client certificates
- This CA is **trusted only inside the organization**.

User/Client gets a certificate:

- Each user or system is given a **client certificate** signed by the internal CA.
- This certificate acts like a **digital ID**.

3 During TLS handshake:

- The **server requests** a client certificate.
- The **client sends** its certificate.
- The server checks the certificate using the **custom CA**.

4 If it's valid, access is granted.

- This proves the user/system is trusted (not just using a password).

 So yes:

“User authentication” happens via certificates in private systems, and a **custom CA** is used to issue and verify those certs.

This makes the system **highly secure**—even if someone knows our password, they **can't connect without the valid certificate**.

Scenario

We are already an **admin** of a Kubernetes cluster.

Now, a **new admin** (say some x person) joins the team.

We want to give them access using **certificate-based authentication**.

Step-by-Step: Client Certificate Authentication for New Kubernetes User

Step 1: New User Generates Private Key and CSR

This step is done **on the new user's local machine** (or securely managed environment).

```
# 1. Generate a private key (3072-bit RSA)
openssl genrsa -out admin.key 3072

# 2. Create a CSR (Certificate Signing Request)
openssl req -new -key admin.key -out admin.csr -
```

- ◆ `CN=admin` → Common Name (the username)
- ◆ `O=system:masters` → Group (must be `system:masters` for full admin privileges)

Step 2: Admin (We) Approve the Request in the Cluster

We now take the `.csr` file from the user and **generate a certificate signed by the Kubernetes cluster CA**.

 Method 1: Using Kubernetes CA Directly

This works **only if we have access to the Kubernetes CA key and cert** (usually available on control plane node under `/etc/kubernetes/pki`):

```
root@kind-control-plane:/# cd /etc/kubernetes/pki/
root@kind-control-plane:/etc/kubernetes/pki# ls -lrt
total 60
-rw----- 1 root root 1675 Jul 17 13:03 ca.key
-rw-r--r-- 1 root root 1107 Jul 17 13:03 ca.crt
-rw----- 1 root root 1679 Jul 17 13:03 apiserver-kubelet-client.key
-rw-r--r-- 1 root root 1176 Jul 17 13:03 apiserver-kubelet-client.crt
-rw----- 1 root root 1675 Jul 17 13:03 front-proxy-ca.key
-rw-r--r-- 1 root root 1123 Jul 17 13:03 front-proxy-ca.crt
-rw----- 1 root root 1675 Jul 17 13:03 front-proxy-client.key
-rw-r--r-- 1 root root 1119 Jul 17 13:03 front-proxy-client.crt
drwxr-xr-x 2 root root 4096 Jul 17 13:03 etcd
-rw----- 1 root root 1679 Jul 17 13:03 apiserver-etcd-client.key
-rw-r--r-- 1 root root 1123 Jul 17 13:03 apiserver-etcd-client.crt
-rw----- 1 root root 451 Jul 17 13:03 sa.pub
-rw----- 1 root root 1675 Jul 17 13:03 sa.key
-rw----- 1 root root 1675 Jul 18 11:33 apiserver.key
-rw-r--r-- 1 root root 1326 Jul 18 11:33 apiserver.crt
root@kind-control-plane:/etc/kubernetes/pki# |
```

```
# Sign the CSR using the Kubernetes CA
openssl x509 -req -in admin.csr -CA /etc/kubernetes/certs/ca.crt
-CAkey /etc/kubernetes/pki/ca.key -CAcreateserial
```

This creates:

- `admin.crt` (user's client certificate)
- `admin.key` (user's private key)

```

Krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ openssl genrsa -out admin.key 3072
Generating RSA private key, 3072 bit long modulus (2 primes)
.....+++++
e is 65537 (0x010001)
Krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ openssl req -new -key admin.key -out admin.csr -subj "/CN=admin"
Krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ cat admin.csr | base64 | tr -d '\n'
LS0tLS1CRUdJTjBDRVJUSUZQOFURSB5RVFVRNULS9tL50K7ULJRFZUQ9NBYjBD0VFBd0VERU9NOxdHOTFVRUF3d0ZZV1J0YVc0d2dnR2lNQTBH01NxR1NJYjNEUUVQVFVQ0pBNElCandBd2dnR0tBb01CZ1FETjQ3b0RhYmVqY0xyZGZoaWfHYtUBeVZC8dGn2Fac5pKXkxZndCn2lQc2xhcnBcHrWe19BKzN2YxkEV3Iw53pfdhZ0b291RFIIYzIq0TdyZGp1eEk1THc1WtL2bFc3y25p00Z1clc1RV/FQ01EkakJ0eU5XNb33UctNemw5S21VtNfCmpCNlhrYTVtbnR1Zhj2TWhB00hE0T
N0L0l4L3l8dTdoBuIsY1ZYazkzZgpSQXubgwdUx0Wll3R0sxrLfzR2swchlvSzBhaWg5eXliWktvToT2m4HJUT2dKz2kS3orNF14YtKVGhJdwFqckVXhENy1hoSv1JbJfQd0UNGTVB2NxPckcmxa2tWUfVkJ0l0WGb8z25QM0FLb1JEK3lwcm0v0tHmF13xXb0cX0b0JMMVu01W0Uf0jFLYm42UEICb0V3z2xdjIvRzBzXWRKtmej1K22jZtmz0endMrjJps593d2VLShngx22UNR01tS3Rt0wxhCmtTczNp0XNb1hvUn2KRk2xdjI4cUhsa0157jNQm3WYXF1Nep1wJjMg2u0elvNG1hckV1vS9KV1ZwUUU1tEkd2tJ03dH20j107JFL1FpTfMnt0ENBd0VB00fBQ01BmEd0U3FH0U1lM0RRRUJD01VBQfTRJ0mdRojJB0N58SE16ZAp
pbTVQqC95MwdnaZFSjMyCwhaRTZSNdZLchREERnwbmI1UV0rYm9YRD1wb2Rr0jdzetZnUe11NVElx0pjRWUzCnJ5YmISYp0NUm40Vb6MG0vVWJSU1s0kveVnzk1hv0bVxUJ2Cr4bnpMblZca21BMPd1ym1oYUVEMtLN20KSmtJL01Jah
HubEZ0N3g3Si3aFpVWk51b0x0afk50U04T2d0o2AxZwpprenZjTmV0nFmYkx2m1GTHkzNUpKRp0dy9T21BrUmE1U3YyK1gx2ZFRXkxK2p1K1U4aEztYT2RTTNT0U1ns1l0cERqawJKUmlq0jZ6dFRGdDJreU12CmpPM0Ezd0Lic3BTWtLZ0
{FzKznh22VH0251Unjt3pviXZJSU23cExXVTFUdy84K3NHTmUyRGJB5Dk2WkdoSDUKQVNGSTBNUl0jjsURyV1lyAfw1WkdpVmhud001NwLyb6l6NwX0ehJzAhHzkwrcnBFR3RRY015MuuaUiZwDJuCwpxamRXV2wNHlDbm1UdkZ4cElNy2kVz
TnBydz1pM0t12EJU0WNRV0yYwhp0V0XnqWVS96QzV0RWh1c050Y2VaCjNfZTh253p10WZob2pVYBKK1FyUDVtcitJQ3NreUfTUlhiKndlaflpWktjZwxFajZ0MnVUZnM9C1otL50tRSEIEfWf1RJRk1D0VRFIFJFUvvfU10tL50tL0o=kru
pakar@krupakar:~/Documents/Docker/k8s/kubernetes$ 
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl apply -f CSR.yaml
certificatesigningrequest.certificates.k8s.io/admin created
Krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl get csr
NAME   AGE   SIGNERNAME           REQUESTOR           REQUESTEDDURATION   CONDITION
admin   12s   kubernetes.io/kube-apiserver-client   kubernetes-admin   24h                Pending
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl certificate approve admin
certificatesigningrequest.certificates.k8s.io/admin approved
Krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl get csr
NAME   AGE   SIGNERNAME           REQUESTOR           REQUESTEDDURATION   CONDITION
admin   17s   kubernetes.io/kube-apiserver-client   kubernetes-admin   24h                Approved,Issued
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl get csr admin -o jsonpath='{.status.certificate}' | base64 --decode > admin.crt
Krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl config set-credentials admin \
>   --client-key=admin.key \
>   --client-certificate=admin.crt \
>   --embed-certs=true
User "admin" set.
Krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl config set-context admin --cluster=kind-kind --user=admin
Context "admin" created.
Krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl config get-contexts
CURRENT   NAME   CLUSTER   AUTHINFO   NAMESPACE
*         admin   kind-kind   admin
*         kind-kind   kind-kind   kind-kind
Krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl config use-context admin
Switched to context "admin".
Krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl config get-contexts
CURRENT   NAME   CLUSTER   AUTHINFO   NAMESPACE
*         admin   kind-kind   admin

```

Extracted the Signed Certificate:

```
kubectl get csr admin -o jsonpath='{.status.cert}
```

This command decodes the signed certificate (`admin.crt`) for the user `admin` after the CSR was approved.

Step 3: Set Up kubeconfig File for the New User

Set User Credentials in Kubeconfig:

Now we configure a `kubeconfig` file for `admin` with their credentials

```
kubectl config set-credentials admin\  
  --client-key=admin.key \  
  --client-certificate=admin.crt \  
  --embed-certs=true
```

- We configured kubeconfig with the user's credentials (private key and signed certificate).
- `--embed-certs=true` embeds the actual certificate and key data inside the kubeconfig file—good for portability.



What's Next? (Step 4: Authorization)

Now that the `admin` user is **authenticated**, Kubernetes recognizes them—but they still **can't access cluster resources** until they are **authorized**.

```
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl auth whoami  
ATTRIBUTE          VALUE  
Username           kubernetes-admin  
Groups             [kubeadm:cluster-admins system:authenticated]  
Extra: authentication.kubernetes.io/credential-id [X509SHA256-ad84cf8d66a8dbabfe1cc4ac191b827e07683a58cb57c57d348f8ce1747a72f]  
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl auth can-i get pods  
yes  
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl auth can-i get pod --as admin  
no  
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ |
```



Authentication vs Authorization in Kubernetes

Authentication = Identity Check

- **In Kubernetes:** The user presents credentials (like a client certificate) to prove who they are.

- **Analogy:** We show our ID card at the office entrance. Security confirms who we are .

Authorization = Access Control

- **In Kubernetes:** Once our identity is confirmed, RBAC rules determine what we're allowed to do.
- **Analogy:** After entering the building, we're told which room/hall to go to (e.g., training room) and which areas we're not allowed to enter (e.g., other departments).

In our example:

- The new employee is authenticated when their identity is verified at the entrance.
- The new employee is authorized when the company decides what rooms or resources they can access.

```
krupakar@krupakar:~/.kube$ cd ..
krupakar@krupakar:~$ cd $HOME/.kube
krupakar@krupakar:~/.kube$ ls -lrt
total 12
drwxr-x--- 4 krupakar krupakar 4096 Apr 30 22:30 cache
-rw----- 1 krupakar krupakar 5602 Jul 17 18:33 config
krupakar@krupakar:~/.kube$ |
```

```
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl config use-context admin
Switched to context "admin".
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl config view
apiVersion: v1
clusters:
- cluster:
  certificate-authority-data: DATA+OMITTED
  server: https://127.0.0.1:34477
  name: kind-kind
contexts:
- context:
  cluster: kind-kind
  user: admin
  name: admin
- context:
  cluster: kind-kind
  user: kind-kind
  name: kind-kind
current-context: admin
kind: Config
preferences: {}
users:
- name: admin
  user:
    client-certificate-data: DATA+OMITTED
    client-key-data: DATA+OMITTED
- name: kind-kind
  user:
    client-certificate-data: DATA+OMITTED
    client-key-data: DATA+OMITTED
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ |
```

The user entry gets recorded in the config file

After the user is authenticated (their identity is verified), Kubernetes needs to **authorize** what they can do—like viewing pods, deleting deployments, etc.

There are 4 main **authorization modes** in Kubernetes:

We should now assign RBAC permissions:

- ◆ **1. RBAC (Role-Based Access Control)**
 - Most commonly used mode.
 - Access is given based on **roles** assigned to **users or groups**.
 - Defined using:
 - `Role` , `ClusterRole`

- `RoleBinding` , `ClusterRoleBinding`

📌 **Use case:** “Give read-only access to a user in a namespace.”

◆ **2. ABAC (Attribute-Based Access Control)**

- Access is granted based on a **policy file** that contains rules with user attributes.
- Not dynamic like RBAC; needs changes in static files.
- Rarely used in production now.

📌 **Use case:** Older systems with complex conditions or where RBAC is not sufficient.

◆ **3. Node Authorization**

- Used **only by kubelets (nodes)**.
- Grants permissions to nodes (e.g., reading secrets or pod info of its own pods).
- Prevents nodes from accessing other node's data.

📌 **Use case:** Allow a node to read its own pod info but **not other nodes' data**.

◆ **4. Webhook Authorization**

- Kubernetes sends a request to an **external service** (like a custom API).
- The external service **decides** whether to allow the request.
- Used in **custom or enterprise** setups.

📌 **Use case:** “Allow access only during office hours”—logic handled outside Kubernetes.

Role, RoleBinding, and ClusterRoleBinding

◆ What is a Role?

A `Role` defines permissions (what actions are allowed) **within a specific namespace**.

- 🔐 It lists what verbs (get, list, create, delete, etc.) a user can perform on which Kubernetes resources (pods, secrets, configmaps, etc.).
- 📂 Namespace-scoped—it only applies to one namespace.

Think of this like assigning permissions **within one department** (e.g., Sales team folder or Development namespace).

```
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl auth whoami
ATTRIBUTE          VALUE
Username           kubernetes-admin
Groups             [kubeadm:cluster-admins system:authenticated]
Extra: authentication.kubernetes.io/credential-id [X509SHA256=ad84cf8d666a8dbabfe1cc4ac191b827e07683a58cb57c57d348f8ce1747a72f]
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl auth can-i get pods
yes
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl auth can-i get pod --as admin
no
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ |
```

```

krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl api-resources --namespaced=true
NAME           SHORTNAMES   APIVERSION          NAMESPACED   KIND
bindings       v1           v1                  true        Binding
configmaps    cm          v1                  true        ConfigMap
endpoints     ep          v1                  true        Endpoints
events        ev          v1                  true        Event
limitranges   limits      v1                  true        LimitRange
persistentvolumeclaims pvc         v1                  true        PersistentVolumeClaim
pods          po          v1                  true        Pod
podtemplates  v1           v1                  true        PodTemplate
replicationcontrollers rc          v1                  true        ReplicationController
resourcequotas quota      v1                  true        ResourceQuota
secrets        v1           v1                  true        Secret
serviceaccounts sa          v1                  true        ServiceAccount
services       svc         v1                  true        Service
controllerrevisions v1           apps/v1            true        ControllerRevision
daemonsets    ds          apps/v1            true        DaemonSet
deployments   deploy      apps/v1            true        Deployment
replicasets   rs          apps/v1            true        ReplicaSet
statefulsets  sts         apps/v1            true        StatefulSet
localsubjectaccessreviews v1           authorization.k8s.io/v1  true        LocalSubjectAccessReview
horizontalpodautoscalers hpa         v1           autoscaling/v2  true        HorizontalPodAutoscaler
cronjobs      cj          batch/v1           true        CronJob
jobs          v1           batch/v1           true        Job
leases        v1           coordination.k8s.io/v1  true        Lease
endpointslices v1           discovery.k8s.io/v1  true        EndpointSlice
events        ev          events.k8s.io/v1      true        Event
ingresses     ing         networking.k8s.io/v1  true        Ingress
networkpolicies netpol     networking.k8s.io/v1  true        NetworkPolicy
poddisruptionbudgets pdb         policy/v1          true        PodDisruptionBudget
rolebindings  v1           rbac.authorization.k8s.io/v1  true        RoleBinding
roles         v1           rbac.authorization.k8s.io/v1  true        Role
csistoragecapacities v1           storage.k8s.io/v1      true        CSISStorageCapacity
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ |

```

Namespace-scoped

```

krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ code .
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ cat role.yaml
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  namespace: default
  name: pod-reader
rules:
- apiGroups: [""] # "" indicates the core API group
  resources: ["pods"]
  verbs: ["get", "watch", "list"]
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl apply -f role.yaml
role.rbac.authorization.k8s.io/pod-reader created
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl get role
NAME      CREATED AT
pod-reader 2025-07-18T14:04:59Z
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl describe role pod-reader
Name:      pod-reader
Labels:    <none>
Annotations: <none>
PolicyRule:
  Resources  Non-Resource URLs  Resource Names  Verbs
  -----      -----           -----           -----
  pods        [ ]                [ ]             [get watch list]
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ |

```

```

krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl get roles
NAME          CREATED AT
pod-reader    2025-07-18T14:04:59Z
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl get roles -A
NAMESPACE      NAME                                CREATED AT
default        pod-reader                         2025-07-18T14:04:59Z
kube-public    kubeadm:bootstrap-signer-clusterinfo 2025-07-17T13:03:37Z
kube-public    system:controller:bootstrap-signer   2025-07-17T13:03:35Z
kube-system    extension-apiserver-authentication-reader 2025-07-17T13:03:34Z
kube-system    kube-proxy                          2025-07-17T13:03:38Z
kube-system    kubeadm:kubelet-config                2025-07-17T13:03:36Z
kube-system    kubeadm:nodes-kubeadm-config         2025-07-17T13:03:36Z
kube-system    system::leader-locking-kube-controller-manager 2025-07-17T13:03:34Z
kube-system    system::leader-locking-kube-scheduler 2025-07-17T13:03:34Z
kube-system    system:controller:bootstrap-signer   2025-07-17T13:03:34Z
kube-system    system:controller:cloud-provider       2025-07-17T13:03:34Z
kube-system    system:controller:token-cleaner      2025-07-17T13:03:34Z
local-path-storage local-path-provisioner-role     2025-07-17T13:03:40Z
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl get roles -A --no-headers
13
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ |

```

◆ What is a RoleBinding?

A `RoleBinding` connects a `user/group/serviceaccount` to a `Role`—essentially saying "this user can have these permissions" in that namespace.

- 🌟 Binds a Role to a user/group **within a namespace**.

This is like saying: "Assign this employee (user) to the Sales department with these specific duties."

```

krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ cat rolebinding.yml
apiVersion: rbac.authorization.k8s.io/v1
# This role binding allows "jane" to read pods in the "default" namespace.
# You need to already have a Role named "pod-reader" in that namespace.
kind: RoleBinding
metadata:
  name: read-pods
  namespace: default
subjects:
# You can specify more than one "subject"
- kind: User
  name: admin # "name" is case sensitive
  apiGroup: rbac.authorization.k8s.io
roleRef:
  # "roleRef" specifies the binding to a Role / ClusterRole
  kind: Role #this must be Role or ClusterRole
  name: pod-reader # this must match the name of the Role or ClusterRole you wish to bind to
  apiGroup: rbac.authorization.k8s.io
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl apply -f rolebinding.yml
rolebinding.rbac.authorization.k8s.io/read-pods created
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl get rolebinding
NAME        ROLE          AGE
read-pods   Role/pod-reader 14s
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl describe rolebinding read-pods
Name:         read-pods
Labels:       <none>
Annotations: <none>
Role:
  Kind:  Role
  Name:  pod-reader
Subjects:
  Kind  Name  Namespace
  ----  ----  -----
  User  admin
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl auth can-i get pod --as admin
yes
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ |

```

◆ What is a ClusterRole?

A `ClusterRole` defines permissions like a Role, but it can be applied cluster-wide (across all namespaces).

- It can also define **non-namespaced resources** (like nodes, persistent volumes).
- 💡 It can be used in two ways:
 1. Grant access to resources across all namespaces.
 2. Grant access to cluster-scoped resources.

Think of this like assigning access to **entire company-wide systems** (like HR software, internal email systems, etc.)

```

krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl auth can-i get nodes
Warning: resource 'nodes' is not namespace scoped

yes
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl auth can-i get nodes --as admin
Warning: resource 'nodes' is not namespace scoped

no
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl create clusterrole node-reader --verb=get,list,watch --resource=node
clusterrole.rbac.authorization.k8s.io/node-reader created
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl get clusterrole
NAME                                CREATED AT
admin                               2025-07-17T13:03:28Z
cluster-admin                       2025-07-17T13:03:28Z
edit                                2025-07-17T13:03:28Z
kindnet                            2025-07-17T13:03:39Z
kubeadm:get-nodes                  2025-07-17T13:03:36Z
local-path-provisioner-role         2025-07-17T13:03:40Z
node-reader                          2025-07-18T17:14:30Z
system:aggregate-to-admin           2025-07-17T13:03:28Z
system:aggregate-to-edit            2025-07-17T13:03:28Z
system:aggregate-to-view            2025-07-17T13:03:28Z
system:auth-delegator              2025-07-17T13:03:28Z
system:basic-user                  2025-07-17T13:03:28Z
system:certificates.k8s.io:certificatesigningrequests:nodeclient   2025-07-17T13:03:29Z
system:certificates.k8s.io:certificatesigninarequests:selfnodeclient    2025-07-17T13:03:29Z

krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl get nodes
NAME      STATUS    ROLES      AGE     VERSION
kind-control-plane  NotReady  control-plane  28h    v1.33.1
kind-worker        Ready     <none>     28h    v1.33.1
kind-worker2       Ready     <none>     28h    v1.33.1
Krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl delete node kind-worker
Error from server (Forbidden): nodes "kind-worker" is forbidden: User "admin" cannot delete resource "nodes" in API group "" at the cluster scope
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ |

```

◆ What is a ClusterRoleBinding?

A `ClusterRoleBinding` connects a user/group/serviceaccount to a `ClusterRole`—giving them access to cluster-level or all namespaces.

Grants cluster-wide access.

This is like giving a new employee company-wide admin access—they can walk into any department, use any system.

Final Flow (from joining to access)

1.  Employee joins company → Gets ID (Create certificate and user in Kubernetes).
2.  Authentication → Security checks ID (Kubernetes verifies client cert).

3.  **Authorization** → Management tells which dept/resource he can access.
4.  **Assign Role** → In dev team, can only read pods.
5.  **RoleBinding** → Actually binds employee to that dev role.
6.  **If Admin** → Gets full ClusterRole + ClusterRoleBinding.

NAME	SHORTNAMES	APIVERSION	NAMESPACED	KIND
componentstatuses	cs	v1	false	ComponentStatus
namespaces	ns	v1	false	Namespace
nodes	no	v1	false	Node
persistentvolumes	pv	v1	false	PersistentVolume
mutatingwebhookconfigurations		admissionregistration.k8s.io/v1	false	MutatingWebhookConfiguration
validatingadmissionpolicies		admissionregistration.k8s.io/v1	false	ValidatingAdmissionPolicy
validatingadmissionpolicybindings		admissionregistration.k8s.io/v1	false	ValidatingAdmissionPolicyBinding
validatingwebhookconfigurations		admissionregistration.k8s.io/v1	false	ValidatingWebhookConfiguration
customresourcedefinitions	crd, crds	apiextensions.k8s.io/v1	false	CustomResourceDefinition
apiservices		apiregistration.k8s.io/v1	false	APIService
selfsubjectreviews		authentication.k8s.io/v1	false	SelfSubjectReview
tokenreviews		authentication.k8s.io/v1	false	TokenReview
selfsubjectaccessreviews		authorization.k8s.io/v1	false	SelfSubjectAccessReview
selfsubjectrulesreviews		authorization.k8s.io/v1	false	SelfSubjectRulesReview
subjectaccessreviews		authorization.k8s.io/v1	false	SubjectAccessReview
certificatesigningrequests	csr	certificates.k8s.io/v1	false	CertificateSigningRequest
flowschemas		flowcontrol.apiserver.k8s.io/v1	false	FlowSchema
prioritylevelconfigurations		flowcontrol.apiserver.k8s.io/v1	false	PriorityLevelConfiguration
ingressclasses		networking.k8s.io/v1	false	IngressClass
ipaddresses	ip	networking.k8s.io/v1	false	IPAddress
servicecidrs		networking.k8s.io/v1	false	ServiceCIDR
runtimeclasses		node.k8s.io/v1	false	RuntimeClass
clusterrolebindings		rbac.authorization.k8s.io/v1	false	ClusterRoleBinding
clusterroles		rbac.authorization.k8s.io/v1	false	ClusterRole
priorityclasses	pc	scheduling.k8s.io/v1	false	PriorityClass
csidrivers		storage.k8s.io/v1	false	CSIDriver
csinodes		storage.k8s.io/v1	false	CSINode
storageclasses	sc	storage.k8s.io/v1	false	StorageClass
volumeattachments		storage.k8s.io/v1	false	VolumeAttachment

cluster-level resources

```

krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl get clusterrole |grep node-reader
          2025-07-18T17:14:30Z
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl describe clusterrole node-reader
Name:      node-reader
Labels:    <none>
Annotations: <none>
PolicyRule:
  Resources  Non-Resource URLs  Resource Names  Verbs
  *          *                  *              *
  nodes      []                [get, list, watch]
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl create clusterrolebinding reader-bind --clusterrole=node-reader --user=admin
error: unknown flag: --clusterrole
See 'kubectl create --help' for usage.
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl create clusterrolebinding reader-bind --clusterrole=node-reader --user=admin
clusterrolebinding.rbac.authorization.k8s.io/reader-bind created
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl get clusterrolebinding |grep reader
  ClusterRole/node-reader
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl describe clusterrolebinding reader-bind
Name:      reader-bind
Labels:    <none>
Annotations: <none>
Role:
  Kind:  ClusterRole
  Name:  node-reader
Subjects:
  Kind  Name  Namespace
  *    *      *
  User  admin
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ |

```

Kubernetes Service Account?

In Kubernetes, **Service Accounts (SAs)** are **non-human user identities** used by applications, bots, or controllers to interact with the cluster.

Unlike regular **user accounts** (used by admins or developers), service accounts are for automation or internal cluster tasks.



Analogy

Imagine a company:

- **Human users** are employees: Admins, Developers, QA, etc.—they log in with their own ID to do things manually.
- **Service accounts** are like automated machines: A CCTV monitoring system, or a daily auto-email bot—they perform tasks on behalf of the company without human input.



When and Why Are Service Accounts Used?

- When **applications running inside the cluster** (like Jenkins, Prometheus, ArgoCD, etc.) need to talk to the Kubernetes API.

- For automation tasks like:
- A Jenkins pipeline deploying apps.
- A monitoring tool scraping metrics.
- A cronjob interacting with the cluster.

Human users **don't trigger these manually**, so we don't want to assign our own credentials. We want a **dedicated identity**: that's the service account.



How Are Service Accounts Managed?

1. Automatically Created:

Every namespace has a default service account.

Run this command:

```
kubectl get sa -A | grep default
```

1. If there are 5 namespaces, we'll see 5 default service accounts (1 per namespace).

2. Usage by Pods:

Every pod automatically uses the **default service account of its namespace**—unless we tell it to use a specific one.

```

krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl get sa
NAME      SECRETS   AGE
default    0          2d2h
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl get sa -A
NAMESPACE     NAME           SECRETS   AGE
default       default        0          2d2h
kube-node-lease default        0          2d2h
kube-public   default        0          2d2h
kube-system   attachdetach-controller 0          2d2h
kube-system   bootstrap-signer    0          2d2h
kube-system   certificate-controller 0          2d2h
kube-system   clusterrole-aggregation-controller 0          2d2h
kube-system   coredns         0          2d2h
kube-system   cronjob-controller 0          2d2h
kube-system   daemon-set-controller 0          2d2h
kube-system   default         0          2d2h
kube-system   deployment-controller 0          2d2h
kube-system   disruption-controller 0          2d2h
kube-system   endpoint-controller 0          2d2h
kube-system   endpointslice-controller 0          2d2h
kube-system   endpointslicemirroring-controller 0          2d2h
kube-system   ephemeral-volume-controller 0          2d2h
kube-system   expand-controller   0          2d2h
kube-system   generic-garbage-collector 0          2d2h
kube-system   horizontal-pod-autoscaler 0          2d2h
kube-system   job-controller    0          2d2h
kube-system   kindnet          0          2d2h
kube-system   kube-proxy        0          2d2h
kube-system   legacy-service-account-token-cleaner 0          2d2h
kube-system   namespace-controller 0          2d2h
kube-system   node-controller   0          2d2h

```

Default



Why Does Every Namespace Have a Default Service Account?

When Kubernetes creates a **new namespace**, it automatically creates a **default service account** inside it.



The Reason: So Pods Can Talk to the Kubernetes API (If Needed)

Here's why:

- When we deploy a Pod, that Pod **might need to communicate with the Kubernetes API server**—for example:
 - To check its own status.
 - To watch for ConfigMap changes.

- To interact with other services.
- Kubernetes needs a way to **authenticate** that Pod to the API server.
- That's where the **default service account** comes in—it's like giving every Pod a badge saying:
"Hi, I'm from namespace X. Here's my identity."
- If we don't explicitly assign a service account in our pod spec, Kubernetes **assigns the default one in that namespace**.

3. Assigning Permissions:

Just like human users, we can assign roles to service accounts using:

- `RoleBinding` (within a namespace)
- `ClusterRoleBinding` (for cluster-wide access)

1. Example:

If Jenkins needs to deploy apps, we create a `jenkins-service-account`, and bind it to a role that gives it the required access (like create pods, manage deployments, etc.).



Analogy

Imagine we join a new department in a company (a namespace in Kubernetes).

- The company automatically gives every department a **generic access badge** (default service account).
- If a machine (pod) is installed in that department, it **gets that default badge** unless we request a specific one.

This ensures the machine has **some identity** if it ever needs to ask for internal resources or perform tasks.

Why Not Use Human Users?

Let's say there are 100 people in the Operations team. If we use their personal accounts to run automation like Jenkins pipelines:

- Logs will show *who* ran it, even if it was triggered by code.
- If a person leaves the company, automation breaks.
- It's insecure, unscalable, and hard to manage.

Service accounts solve all this: consistent identity, proper RBAC, easily manageable.

Use Case:

When a **pod** runs, it pulls the **container image** from a registry (like Docker Hub or ECR).

If the image is in a **public registry**, no login is needed.

But if the image is in a **private registry**, Kubernetes needs **credentials** to pull that image.

Solution: Image Pull Secret

Kubernetes uses an `imagePullSecret` to store Docker registry **credentials** securely.

Then, it can use that secret to log in and **pull the private image**.



How to create an image pull secret?

Let's say our Docker registry username is `myuser` and password is `mypassword`.

```
kubectl create secret docker-registry my-registry  
  --docker-username=myuser \  
  --docker-password=mypassword \  
  --docker-email=myemail@example.com \  
  --docker-server=https://index.docker.io/v1/
```

This creates a secret of type `kubernetes.io/dockerconfigjson`.

🚀 How to use it in a Pod?

In the pod YAML:

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: private-image-pod  
spec:  
  containers:  
    - name: myapp  
      image: myuser/myprivateimage:latest  
  imagePullSecrets:  
    - name: my-registry-secret
```

This tells the pod to use `my-registry-secret` to log in and pull the image.

```
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl auth can-i get pods --as test-sa
no
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl create role test-role \
> --verb=list,get,watch \
> resource=pod
error: exactly one NAME is required, got 2
See 'kubectl create role -h' for help and examples
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl create role test-role --verb=list,get,watch --resource=pod
role.rbac.authorization.k8s.io/test-role created
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl create rolebinding test-rb --role=test-role --user=test-sa
rolebinding.rbac.authorization.k8s.io/test-rb created
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl get pods --as test-sa
No resources found in default namespace.
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl auth can-i get pods --as test-sa
yes
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ |
```



Do We need to add it for every pod?

Yes—by default, we have to specify `imagePullSecrets` for each pod using a private image.

But we can attach this once to a `ServiceAccount`, and all pods using that SA will use the secret.

```

krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl get pods --as test-sa
NAME    READY   STATUS    RESTARTS   AGE
nginx  1/1     Running   0          14s
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl describe pod nginx
Error from server (NotFound): pods "nginx" not found
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl describe pod nginx
Name:           nginx
Namespace:      default
Priority:       0
Service Account: default
Node:           kind-worker2/172.18.0.3
Start Time:     Sat, 19 Jul 2025 21:55:37 +0530
Labels:         run=nginx
Annotations:    <none>
Status:         Running
IP:            10.244.1.3
IPs:
  IP:  10.244.1.3
Containers:
  nginx:
    Container ID:  containerd://b982468db94dc6f83cd17a541c2935ca9a3815793fbec84e84a964d37496ea86
    Image:          nginx
    Image ID:      docker.io/library/nginx@sha256:f5c017fb33c6db484545793ffb67db51cdd7daebee472104612f73a85063f889
    Port:          <none>
    Host Port:    <none>
    State:         Running
      Started:    Sat, 19 Jul 2025 21:55:43 +0530
    Ready:         True
    Restart Count: 0
    Environment:   <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-dpqf9 (ro)
Conditions:
  Type        Status
  PodReadyToStartContainers  True
  Initialized  True
  Ready        True
  ContainersReady  True
  PodScheduled  True
Volumes:
  kube-api-access-dpqf9:
    Type:          Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName:   kube-root-ca.crt
    Optional:       false
    DownwardAPI:    true

```

Attach Image Pull Secret to Service Account

```
kubectl patch serviceaccount default \
-p '{"imagePullSecrets": [{"name": "my-registr' 
```

This adds the image pull secret to the **default service account** in the namespace.

So now, **any pod** using the default SA will use the secret automatically.

Where is the token stored?

When a pod is created using a service account, the **API token** (used to talk to Kubernetes API) is auto-mounted at:

```
/var/run/secrets/kubernetes.io/serviceaccount/token
```

But this is **different from image pull secret**.

This is for **API access**, while `imagePullSecret` is for **pulling container images**.

In Simple Steps

1. We have a **private image** in Docker Hub (or another registry).
2. We create a **Docker registry secret** in Kubernetes.
3. We either.
 - Add the secret in our pod YAML using `imagePullSecrets` , or
 - Attach it once to the service account (like `default`) so all pods can use it.
4. Kubernetes uses that secret to **login to the registry** and **pull the image**.

Network Policy (Kubernetes)

A **NetworkPolicy** in Kubernetes is a resource that controls how **pods** communicate with each other and with the **outside world**.

It helps allow or block traffic between pods based on rules like labels, namespaces, ports, and IP blocks.



1. Application Architecture Example

Imagine a typical 3-tier app deployed in Kubernetes:

- Frontend Pod (Web) → Port 80
- Backend Pod (App logic) → Port 443
- Database Pod (DB) → Port 3306

The flow of a user request looks like this:

User → Frontend → Backend → Database

This is the **normal communication flow**—it happens in two directions:



2. Types of Network Flows

1. Ingress

Traffic coming into a pod from outside (e.g., User to Frontend)

2. Egress

Traffic going out of a pod (e.g., Backend calling the Database)



3. Default Kubernetes Behavior

By default:



All pods can talk to all other pods in the cluster!

- Frontend can talk to Backend ✓
- Backend can talk to Database ✓

- But **Frontend can also talk directly to Database**
Traffic going out of a pod (e.g., Backend calling the Database)
 3. Default Kubernetes Behavior
By default:
 -  All pods can talk to all other pods in the cluster!
Frontend can talk to Backend 
Backend can talk to Database 
But Frontend can also talk direc
 - **Database can also reply directly to Frontend**

This is **not secure**—it breaks the proper layered access.



4. Why We Need Network Policies

We want to **control** which pods can talk to which—for security and proper design.

So we write **Network Policies** like:

- Frontend **cannot** access Database
- Frontend can only talk to Backend
- Backend can talk to DB
- DB cannot talk to Frontend directly



5. CNI Plugins and Policy Enforcement

- Kubernetes needs a **CNI plugin** that supports Network Policies.
- **Kind** uses `kindnet` by default—it does **not support** enforcing network policies.

```

krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl get po -n=kube-system
NAME                               READY   STATUS    RESTARTS   AGE
coredns-674b8bbfcf-skzt7          1/1    Running   2 (57m ago)  45h
coredns-674b8bbfcf-sv9xh          1/1    Terminating   0          2d19h
coredns-674b8bbfcf-tx457          1/1    Terminating   0          2d19h
coredns-674b8bbfcf-v7vxq          1/1    Running   2 (57m ago)  45h
etcd-kind-control-plane           1/1    Running   0          2d19h
kindnet-bnx66                      1/1    Running   3 (57m ago)  2d19h
kindnet-dcd6n                      1/1    Running   3 (57m ago)  2d19h
kindnet-flzm2                      1/1    Running   0          2d19h
kube-apiserver-kind-control-plane 1/1    Running   0          2d19h
kube-controller-manager-kind-control-plane 1/1    Running   0          2d19h
kube-proxy-g298v                  1/1    Running   3 (57m ago)  2d19h
kube-proxy-j78rm                  1/1    Running   0          2d19h
kube-proxy-vm7wt                  1/1    Running   3 (57m ago)  2d19h
kube-scheduler-kind-control-plane 1/1    Running   0          2d19h
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl get ds -A
NAMESPACE   NAME            DESIRED   CURRENT   READY   UP-TO-DATE   AVAILABLE   NODE SELECTOR   AGE
kube-system kindnet         3         3         2        3           2           kubernetes.io/os=linux   2d19h
kube-system kube-proxy      3         3         2        3           2           kubernetes.io/os=linux   2d19h
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ 

```

- So we install Calico (or Cilium) to support and enforce network policies in Kind.

```

krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ cat kindnetwork.yml
kind: Cluster
apiVersion: kind.x-k8s.io/v1alpha4
nodes:
- role: control-plane
  extraPortMappings:
  - containerPort: 30001
    hostPort: 30001
- role: worker
- role: worker
networking:
  disableDefaultCNI: true
  podSubnet: 192.168.0.0/16
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kind create cluster --name kind-kind --config kindnetwork.yml
Creating cluster "kind-kind" ...
✓ Ensuring node image (kindest/node:v1.33.1)
✓ Preparing nodes 🐳 🐳 🐳
✓ Writing configuration
✓ Starting control-plane 🚀
✓ Installing StorageClass
✓ Joining worker nodes 🐳
Set kubectl context to "kind-kind-kind"
You can now use your cluster with:

kubectl cluster-info --context kind-kind-kind

Thanks for using kind! 😊
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl get nodes
NAME             STATUS   ROLES      AGE   VERSION
Kind-kind-control-plane  NotReady control-plane 4m5s   v1.33.1
Kind-kind-worker   NotReady <none>     86s   v1.33.1
Kind-kind-worker2  NotReady <none>     86s   v1.33.1
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl apply -f https://raw.githubusercontent.com/projectcalico/calico/v3.30.2/manifests/calico.yaml
poddisruptionbudget.policy/calico-kube-controllers created
serviceaccount/calico-kube-controllers created
serviceaccount/calico-node created
serviceaccount/calico-cni-plugin created

```

```

Every 2.0s: kubectl get pods -l k8s-app=calico-node -A

```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	calico-node-5rzp7	0/1	Init:0/3	0	92s
kube-system	calico-node-ls842	0/1	Init:0/3	0	93s
kube-system	calico-node-t55fx	0/1	Init:0/3	0	92s

🔧 Scenario: NetworkPolicy to Control Access Between Frontend → Backend → DB

We'll do this in a structured way:

1 Create 3 Deployments + 3 Services

- frontend
- backend
- database

Each Deployment will:

- Have labels like `role: front` , `role: back` , `role: db`
- Have its own ClusterIP service

```
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl apply -f NWfe.yml
pod/frontend created
service/frontend created
pod/backend created
service/backend created
service/db created
pod/mysql created
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl get po -o wide
NAME      READY   STATUS    RESTARTS   AGE     IP           NODE      NOMINATED-NODE   READINESS   GATES
backend   0/1     ContainerCreating   0        32s    <none>       kind-kind-worker2   <none>      <none>
frontend  0/1     ContainerCreating   0        41s    <none>       kind-kind-worker2   <none>      <none>
mysql     0/1     ContainerCreating   0        16s    <none>       kind-kind-worker2   <none>      <none>
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl get svc -o wide
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE     SELECTOR
backend   ClusterIP  10.96.120.81   <none>        80/TCP      38s    role=backend
db        ClusterIP  10.96.211.213  <none>        3306/TCP   30s    name=mysql
frontend  ClusterIP  10.96.114.128  <none>        80/TCP      46s    role=frontend
kubernetes ClusterIP  10.96.0.1     <none>        443/TCP    9m13s   <none>
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl get po -o wide --watch
NAME      READY   STATUS    RESTARTS   AGE     IP           NODE      NOMINATED-NODE   READINESS   GATES
backend   0/1     ContainerCreating   0        54s    <none>       kind-kind-worker2   <none>      <none>
frontend  0/1     ContainerCreating   0        63s    <none>       kind-kind-worker2   <none>      <none>
mysql     0/1     ContainerCreating   0        359s   <none>       kind-kind-worker2   <none>      <none>
^C
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl get po -o wide --watch
NAME      READY   STATUS    RESTARTS   AGE     IP           NODE      NOMINATED-NODE   READINESS   GATES
backend   0/1     ContainerCreating   0        2m50s   <none>       kind-kind-worker2   <none>      <none>
frontend  0/1     ContainerCreating   0        2m59s   <none>       kind-kind-worker2   <none>      <none>
mysql     0/1     ContainerCreating   0        2m34s   <none>       kind-kind-worker2   <none>      <none>
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl get po -o wide --watch
NAME      READY   STATUS    RESTARTS   AGE     IP           NODE      NOMINATED-NODE   READINESS   GATES
backend   0/1     ContainerCreating   0        3m      <none>       kind-kind-worker2   <none>      <none>
frontend  0/1     ContainerCreating   0        3m9s    <none>       kind-kind-worker2   <none>      <none>
mysql     0/1     ContainerCreating   0        2m44s   <none>       kind-kind-worker2   <none>      <none>
mysql     0/1     ContainerCreating   0        3m10s   <none>       kind-kind-worker2   <none>      <none>
backend   0/1     ContainerCreating   0        3m26s   <none>       kind-kind-worker2   <none>      <none>
frontend  0/1     ContainerCreating   0        3m38s   <none>       kind-kind-worker2   <none>      <none>
mysql     1/1     Running          0        6m5s    192.168.197.2   kind-kind-worker2   <none>      <none>
backend   1/1     Running          0        8m11s   192.168.197.1   kind-kind-worker2   <none>      <none>
frontend  1/1     Running          0        8m34s   192.168.197.3   kind-kind-worker2   <none>      <none>
```

2 Default Behavior (No Network Policy)

- We `kubectl exec` into any pod (e.g., frontend pod)
- We can `curl` all services (backend, db)

- That means: all pods can access all pods

```
curl backend → OK
curl db → OK
curl frontend from db → OK
```

 This is insecure

```
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl exec -it frontend -- sh
# curl backend:80
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
html { color-scheme: light dark; }
body { width: 35em; margin: 0 auto;
font-family: Tahoma, Verdana, Arial, sans-serif; }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>

<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>

<p><em>Thank you for using nginx.</em></p>
</body>
</html>
# telnet db 3306
Trying 10.96.211.213...
|
```

```
# telnet backend 80
Trying 10.96.120.81...
Connected to backend.
Escape character is '^>'.
*Connection closed by foreign host.
# telnet db 3306
Trying 10.96.211.213...
telnet: Unable to connect to remote host: Connection refused
# telnet db 3306
Trying 10.96.211.213...
Connected to db.
Escape character is '^>'.
I
9.3.0  =>2K:06Pw_I^*[ZJA caching_sha2_password#08501Got timeout reading communication packetsConnection closed by foreign host.
# |
```

```
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl exec -it backend -- sh
# apt-get update && apt-install telnet -y
Get:1 http://deb.debian.org/debian bookworm InRelease [151 kB]
Get:2 http://deb.debian.org/debian bookworm-updates InRelease [55.4 kB]
Get:3 http://deb.debian.org/debian-security bookworm-security InRelease [48.0 kB]
Get:4 http://deb.debian.org/debian bookworm/main amd64 Packages [8793 kB]
Get:5 http://deb.debian.org/debian bookworm-updates/main amd64 Packages [756 B]
Get:6 http://deb.debian.org/debian-security bookworm-security/main amd64 Packages [272 kB]
Fetched 9320 kB in 9s (1023 kB/s)
Reading package lists... Done
sh: 1: apt-install: not found
# curl fronten:80
curl: (6) Could not resolve host: fronten
# curl frontend:80
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
html { color-scheme: light dark; }
body { width: 35em; margin: 0 auto;
font-family: Tahoma, Verdana, Arial, sans-serif; }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>
<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>
<p><em>Thank you for using nginx.</em></p>
</body>
</html>
# |
```

3 Use `telnet` Instead of `curl` for DB Testing

Since DB doesn't give response on `curl`, we'll use `telnet`:

```
apt update && apt install -y telnet
telnet db 3306
```

4 Create Initial Network Policy

Now we want to restrict the access:

- Frontend can talk to Backend
- Backend can talk to Database
- Frontend should NOT talk to Database

- Database should NOT talk to Frontend

5 How to Write the Network Policy

We'll write a **NetworkPolicy** that allows **Ingress** traffic only from the pod that has `role=front`.

Example:

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-front-to-back
spec:
  podSelector:
    matchLabels:
      role: back
  policyTypes:
  - Ingress
  ingress:
  - from:
    - podSelector:
        matchLabels:
          role: front
```

- This means: **Only frontend pods can talk to backend**
- All others (e.g., DB or other) → **X** denied

Repeat similar for backend → db

6 Add Egress Control (Optional)

We can also control who a pod **can talk to (egress)**

Example: Backend should **only** be allowed to talk to DB:

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: back-egress-to-db
spec:
  podSelector:
    matchLabels:
      role: back
  policyTypes:
  - Egress
  egress:
  - to:
    - podSelector:
        matchLabels:
          role: db
```

```

krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ cat networkpolicy.yml
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: db-test
spec:
  podSelector:
    matchLabels:
      name: mysql
  policyTypes:
  - Ingress
  ingress:
  - from:
    - podSelector:
        matchLabels:
          role: backend
    ports:
    - port: 3306
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl apply -f networkpolicy.yml
networkpolicy.networking.k8s.io/db-test created
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl get netpole
error: the server doesn't have a resource type "netpole"
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl get netpol
NAME   POD-SELECTOR   AGE
db-test   name=mysql   35s
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl describe netpol db-test
Name:         db-test
Namespace:    default
Created on:   2025-07-20 16:35:06 +0530 IST
Labels:       <none>
Annotations: <none>
Spec:
  PodSelector:  name=mysql
  Allowing ingress traffic:
    To Port: 3306/TCP
    From:
      PodSelector: role=backend
  Not affecting egress traffic
  Policy Types: Ingress
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ |

```

7 Test the Policy

Now test again:

- ✓ curl backend from frontend → OK
- curl db from frontend → Denied
- telnet db 3306 from backend → OK
- curl frontend from db → Denied

```

krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl exec -it backend -- sh
# telnet db 3306
Trying 10.96.211.213...
Connected to db.
Escape character is '^]'.
I
9.3.0
gRL`'*C00<RP\W!lpbcaching_sha2_password
|
```

```
krupakar@krupakar:~/Documents/Docker/k8s/kubernetes$ kubectl exec -it frontend -- sh
# curl backend:80
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
html { color-scheme: light dark; }
body { width: 35em; margin: 0 auto;
font-family: Tahoma, Verdana, Arial, sans-serif; }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>

<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>

<p><em>Thank you for using nginx.</em></p>
</body>
</html>
# telnet db 3306
Trying 10.96.211.213...
```

✓ Final Setup Flow:

Frontend → Backend → DB



Visual Guide: How TLS, mTLS, and Certificates Work in Kubernetes

A simplified look at how Kubernetes uses certificates for secure communication and authentication between components.

Traditional TLS vs mTLS Connections

Traditional TLS Encryption



Mutual TLS Encryption



Traditional TLS vs Mutual TLS (mTLS)

The diagram shows the key difference between traditional TLS and mutual TLS:

- In **Traditional TLS**, only the **server presents a certificate**. It's like when we visit a website—our browser checks if the site is secure, but the site doesn't check who we are.

Certificate Viewer: medium.com

X

General Details

Issued To

Common Name (CN) medium.com
Organisation (O) <Not part of certificate>
Organisational Unit (OU) <Not part of certificate>

Issued By

Common Name (CN) WE1
Organisation (O) Google Trust Services
Organisational Unit (OU) <Not part of certificate>

Validity Period

Issued On Sunday 20 July 2025 at 00:29:27
Expires On Saturday 18 October 2025 at 01:29:23

SHA-256 Fingerprints

Certificate e20401d82780ee9d82c85b2f78a74e04a8f8067e2ae0ce3a1848eba15bb4e8cc
Public key e25ab18fef7ca18876a61cd46b202c9c4bec32c0a8f1a0e545c4a7b4e385eeee

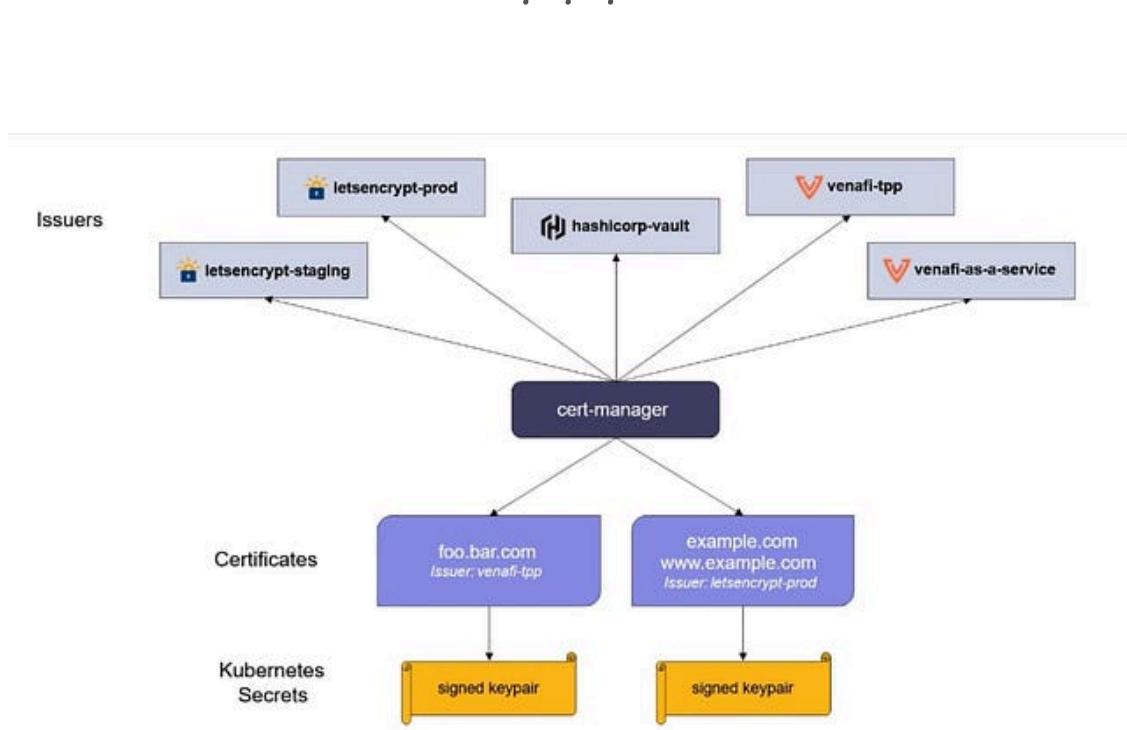
Secure Connection Proof:

Our browser shows this like to confirm the site is real and the connection is safe.

- In **Mutual TLS (mTLS)**, both the client and the server present certificates. It's like a secure handshake where both parties verify each other's identity.

 This extra layer of trust is especially important in Kubernetes environments—like when two internal services (e.g., frontend and

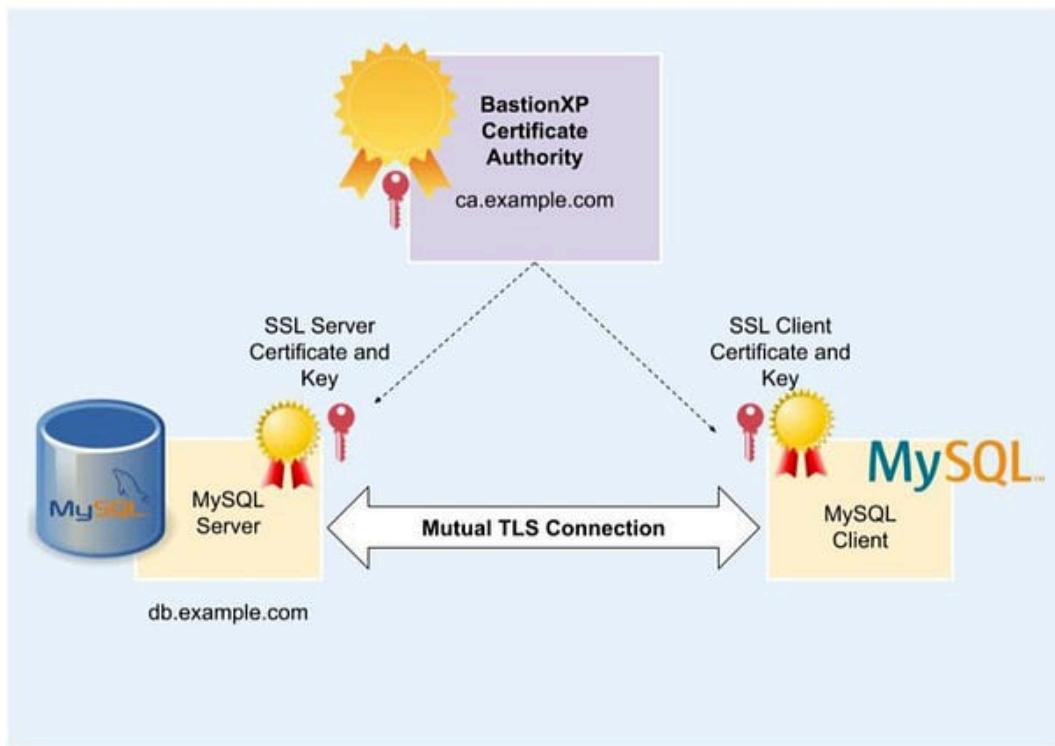
backend) communicate. mTLS ensures that **only trusted components** talk to each other.



The diagram shows how **cert-manager** helps handle certificates automatically inside Kubernetes.

It works with tools like **Let's Encrypt**, **Vault**, or **Venafi** (called *Issuers*) to request and renew certificates for us.

These certificates are saved as Kubernetes **secrets** and used to secure communication using **TLS** or **mTLS**.



Mutual TLS Authentication—MySQL Example

The diagram shows how **mutual TLS (mTLS)** works between a MySQL client and server.

Both the client and server get certificates from the same **Certificate Authority (CA)**.

Before connecting, they **check each other's certificates** to make sure they're trusted.

This way, **both sides are verified**, making the connection secure and trustworthy.

<https://medium.com/@shirishareddy.yasa>